

Reordering sparse matrices for fast multiplication on tensor architectures

Github repository:


<https://github.com/LACSFUB/SPARTA>

Paolo Sylos Labini
Faculty of Computer Science,
Free University of Bozen-Bolzano

Motivation

Tensor cores are good at multiplying dense arrays

But what about **sparse matrices** multiplication?




Should we use sparse-specific data structures (CSR) and algorithms (cusparsparse)?

We can't benefit of the tensor architecture



Should we treat the sparse matrix as a dense one and use tensor architectures?

We end up multiplying a lot of zeros


Our solution



ORIGINAL



REORDERED



BLOCKED

Applications

Reordering takes time

If the same matrix is multiplied repeatedly, we can absorb the cost


- graphs
- neural networks
- matrix collections

The reordering algorithm

Find similar rows and group them together in blocks

- 1) Partitioning the columns
- 2) Grouping by hash
- 3) Grouping by similarity


The reordering algorithm – partitioning the columns



Choice of Δ may depend on

- architecture
- previous knowledge of the matrix
- ...

The reordering algorithm - hashing the rows




1


3

$$h = 1 + 3 = 4$$


The reordering algorithm - hashing the rows



The reordering algorithm – grouping by hash




The reordering algorithm - grouping by hash



- 1) **hashing:** $O(N \times M)$
for CSR matrix with K entries: $O(K)$
- 1) **sorting:** $O(N \log N)$
- 1) **grouping:** approx. $N \times M$

The reordering algorithm - grouping by cosine



For each pair of groups A, B

if $\cos \theta(A, B) \geq \epsilon$

merge A and B


A

| | | |
|---|---|---|
| 0 | 0 | 1 |
|---|---|---|

B

| | | |
|---|---|---|
| 3 | 0 | 3 |
|---|---|---|

The reordering algorithm - grouping by cosine



For each pair of groups A, B


if $\cos \theta(A, B) \geq \epsilon$

merge A and B

The reordering algorithm - grouping by cosine

M

N




$O(N*N*M)$

But with some optimization:


- Only groups are compared, not rows
- Only patterns are compared, not entries

Still, too expensive!

Parallel multiplication



Parallel multiplication




Reorder - experiments with synthetic matrices

Synthetic block matrix


$$\Delta = 2$$

$$\mu = 75\%$$


$$\eta = 50\%$$



Scrambled



Partitioned



$$\Delta' = 3$$

Δ = original block size

μ = original in-block density

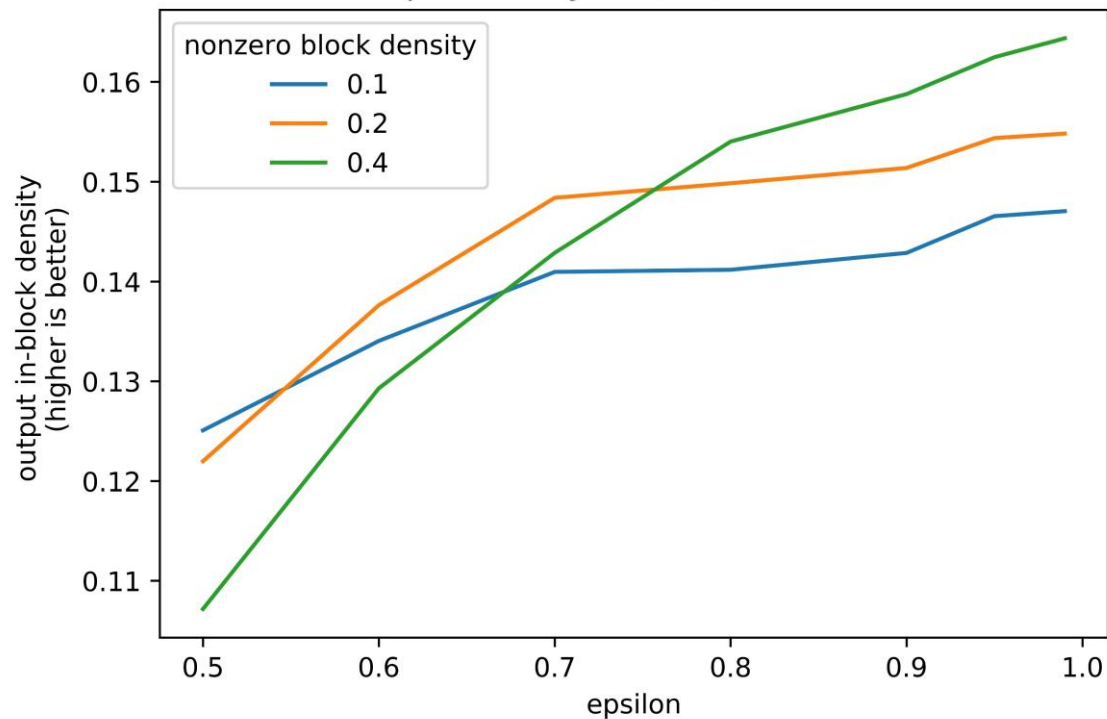
η = original nonzero block density

Reorder - experiments with synthetic matrices

algorithm block size = 35

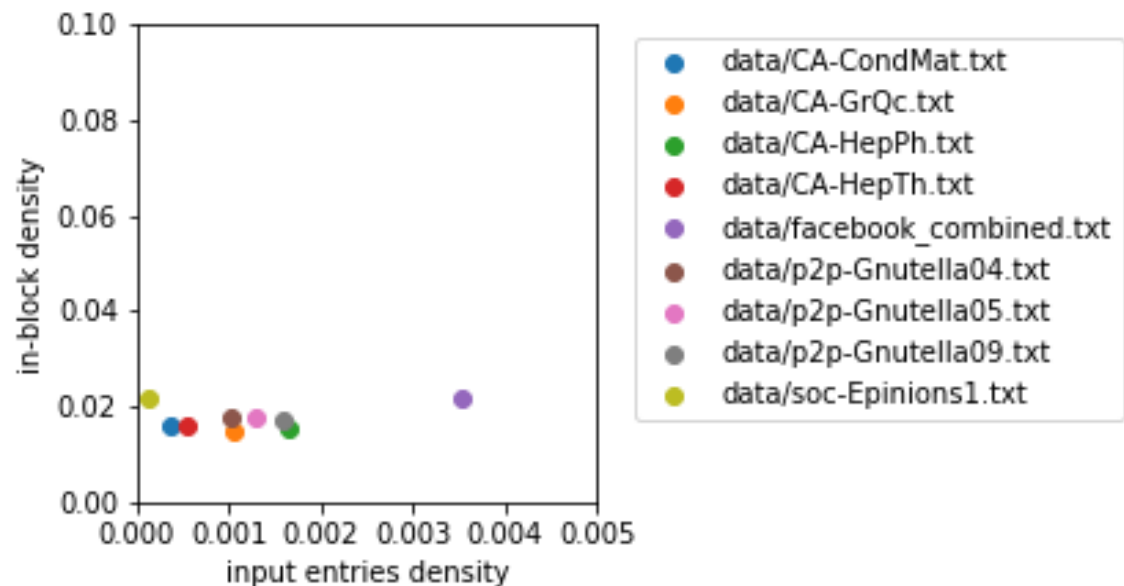
input block size = 64

input density inside blocks = 0.2



Reorder - experiments with real matrices


algo_block_size:64
scramble:1
epsilon:0.6



Reorder - experiments with real matrices


scramble:1

input_source:data/facebook_combined.txt

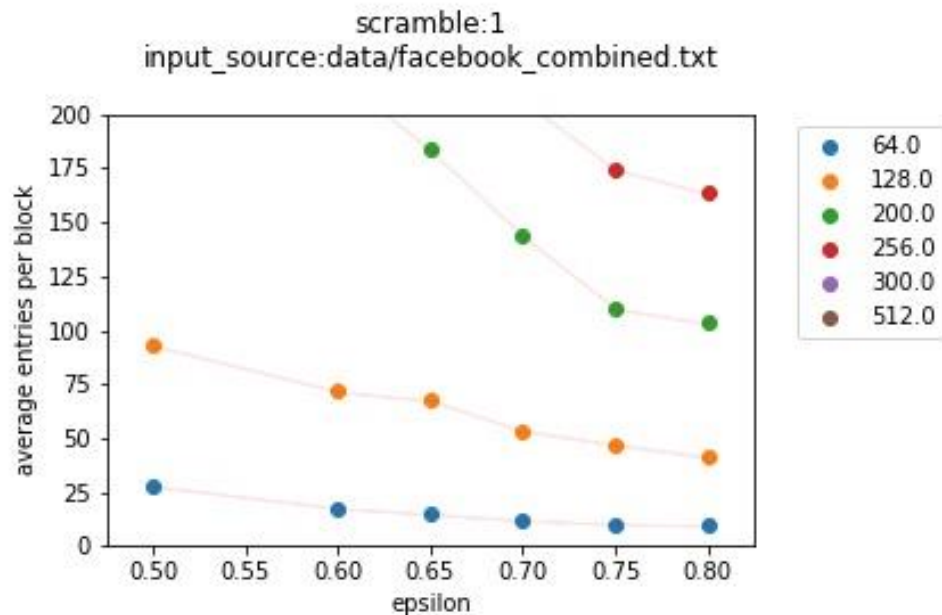
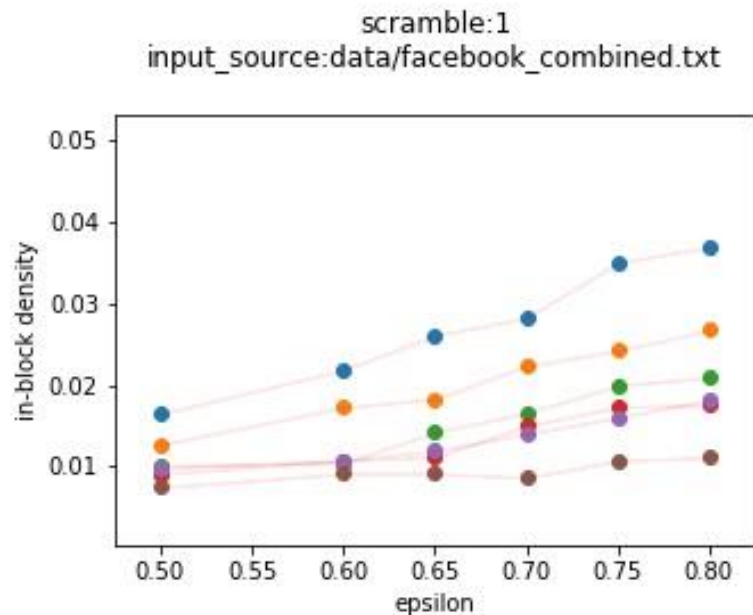


scramble:1

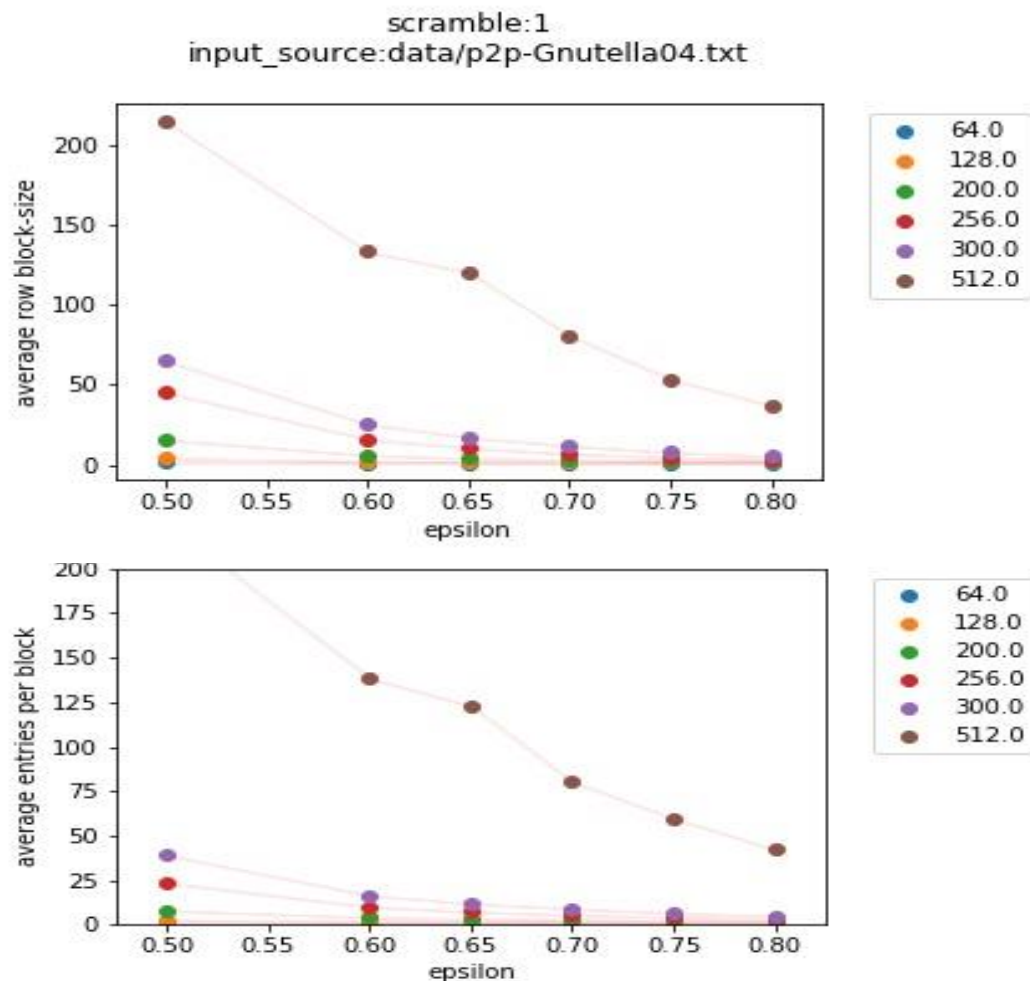
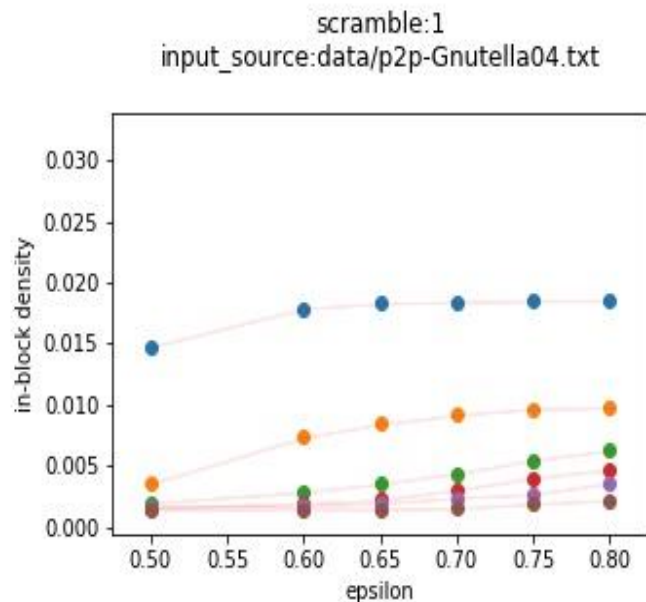
input_source:data/facebook_combined.txt



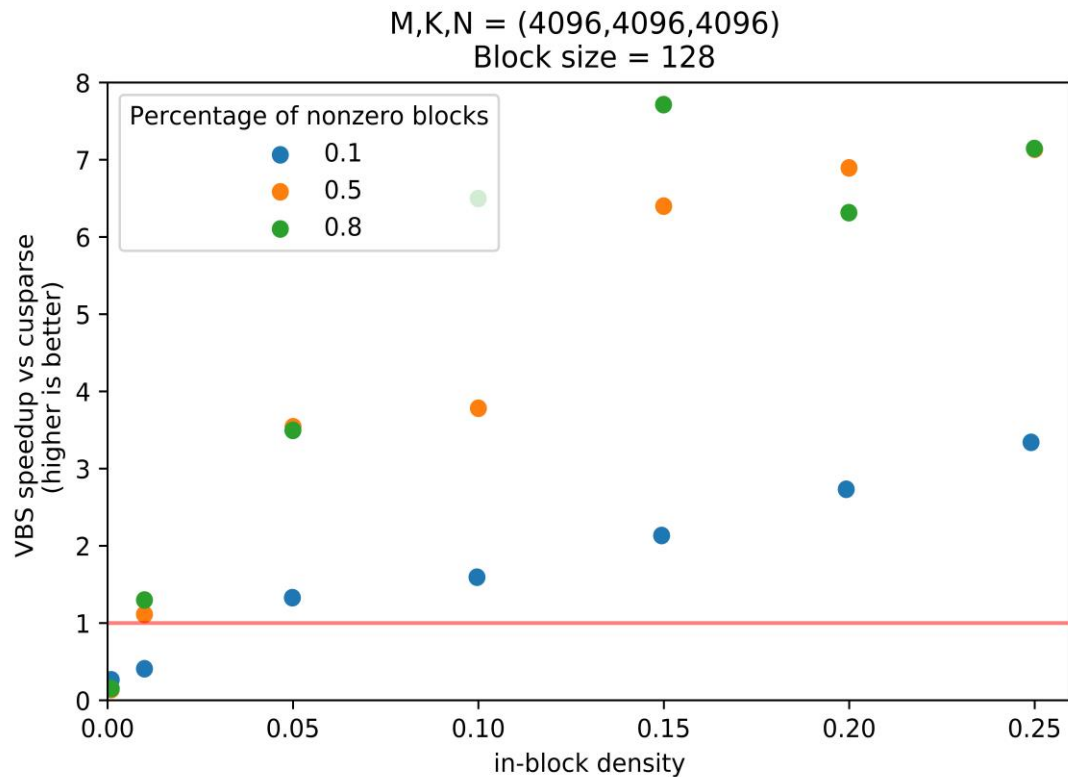
Reorder - experiments with real matrices




Reorder



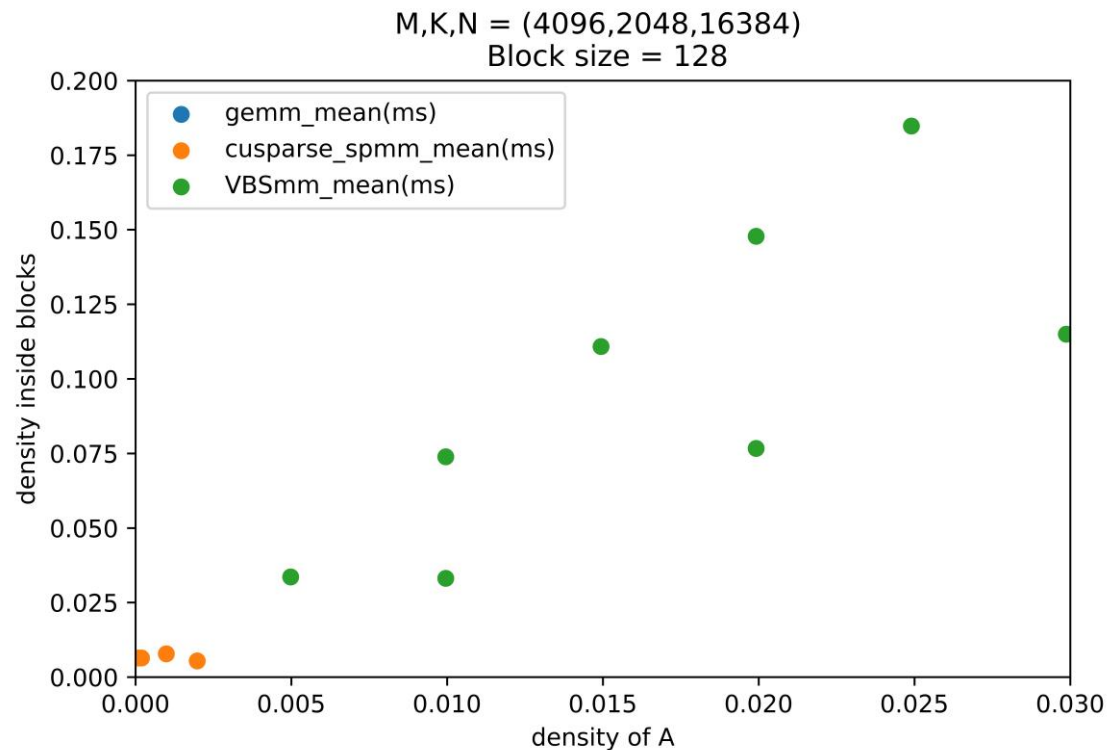
Multiplication - experiments




Multiplication - experiments



Multiplication - experiments



Multiplication - experiments



Future directions – Better, faster reordering

- Parallel implementation
- Local Sensitive Hashing

Future directions – LSH

Faster reordering algorithm through local sensitive hashing and nearest neighbours search


MinHash (random permutations): approximate Jaccard similarity

SimHash (random projections): approximate cosine similarity

Future directions - Compare to related work

Accelerating sparse matrix-matrix multiplication with GPU Tensor
Cores[☆]


Orestis Zachariadis^{a,*}, Nitin Satpute^a, Juan Gómez-Luna^b, Joaquín



Future directions - Tensor-core implementation

- Optimal size of block partition?
- Best way to parallelize?
- Focused experiments

Future directions - Sparse-Sparse multiplication



Problem: zeros gets multiplied by each other
Number of unnecessary operations is squared